


Pseudo-Randomness and Hashing

Pseudo-random number generators

- In previous class we assumed that a randomized algorithm has access to a sequence of IID RV.
- In reality, the program calls a pseudo-random function such as **RAND()**.
- For the sake of simplicity lets assume that each call to Rand generates a single bit.
- The function **RAND** has a persistent d-bit variable called state and it operates in two modes:
 - **RAND(seed)** :
 - Set State=seed.
 - Return $F(\text{State})$ # F returns a bit.
 - **RAND()** :
 - Update $\text{State} = G(\text{State})$ # G returns a new state.
 - Return $F(\text{State})$

A simple random number generator: Linear Feedback Shift Register (LFSR)



- **State**: 16bit number.
- **G**: At each step the bits shift one step to the right and the least significant bit is replaced by the output of the circuit.
- **F**: output the most significant bit (bit 16) as the output of the random number generator;

What does pseudo-random mean?

- We assume that F and G are public knowledge.
- If we know the seed, then we know exactly what the sequence would be.
- Recall what it means that X_1, X_2, \dots, X_n are IID Binary RV with $p = \frac{1}{2}$:
 - $P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n) = \frac{1}{2^n}$
 - $P(X_t = x_t | X_1 = x_1, \dots, X_{t-1} = x_{t-1}) = 1/2$
- In words: it is impossible to predict the bit value of X_t from the values of the sequence so far: X_1, \dots, X_{t-1}

Is it possible to predict the output of a pseudo-random number generator?

- The sequence is determined by the seed.
- If the state has d bits, there are at most 2^d possible seeds and therefore 2^d possible sequences.
- A brute force prediction algorithm: make a list of all possible seeds and the corresponding sequences, after each bit is revealed: delete the seeds whose sequences are inconsistent with the bit.
- It is not hard to show that predicting with the majority of surviving seeds will make no more than d mistakes.
- So pseudo random number generators can be predicted!
- Why do we call them (pseudo) random?
- Because the process of predicting them requires compute resources that are $\Omega(2^d)$.
- In other words, it is computational complexity that gives us pseudo-random number generators!


Pseudo-random vs random

- Random sequence: cannot be predicted.
- Pseudo-Random sequence: Takes prohibitive memory and time to predict.
- LFSR: a very weak pseudo-random number generator: sufficient for applications in signal processing.
- Cryptographically secure pseudo-random number generators: generators for which difficulty of prediction is mathematically proven.
 - Used in cryptographic systems such as RSA, OpenSSL,...
 - Vulnerabilities sometimes related to problems with the pseudo-random generators.

Hash Functions and Hash Tables

Hashing

- A collection of n complex items from a very large set A (names, sounds, images)




Hash Table: Array of pointers of size m

Hash Functions

- Hash functions map items from a very large space (images) to the number $1, \dots, n$
- We want the mapping to behave like a random function: each item $a \in A$ is mapped to $H(a) = i \in \{1, \dots, n\}$ with probability $1/n$ independently of $H(b), H(c), \dots$
- But, we also want the mapping not to be random in that $H(a)$ always gives the same number.
- Solution: instead of one function, we use a family of hash functions, indexed by i : $H_i(a)$ at the start of the program we choose i at random to be a d -bit integer.
- We use a pseudo-random number generator **RAND()** to construct the hash function (assume **RAND()** outputs numbers in the range $1..n$)
 - $H_i(a) = \text{Rand}(i \boxplus a)$ -- $i \boxplus a$ is the binary number created by concatenating the bits of the index i with the bits of the item a

Linked list Hash

- In the single occupancy hash table we discussed last class, each bin in the hash table contains (the pointer to) at most one item.
 - Inserting a new items requires finding an empty bin.
- In the a linked list hash table each bin contains a pointer to a linked list of items
- Unlike single occupancy hash tables, a new item is added to the end of the list starting at the bin to which it is hashed.
- One can store more items than the number of bins in the table.
- The time to add or fetch an item is proportional to the length of the list.



Expected number of items in each bin

m = number of bins n = number of items

X_i = number of items in bin i

Note: $X_1 + X_2 + \cdots + X_m = n$

$$E(X_1 + X_2 + \cdots + X_m) = n$$

What is $E(X_i)$?

A. $E(X_i) = n / n = 1$

B. $E(X_i) = m / m = 1$

C. $E(X_i) = n / m$

C. $E(X_i) = n^2 / m$

Expected running time vs. worst case running time

- As the expected occupancy is n/m and we expect $m > n$, we define $c = \frac{n}{m}$
- Note that any hashing function achieves expected occupancy c . Even one which maps all items to the same bin.
- Beyond expected occupancy, we would like a guarantee that the maximal occupancy is small.
- An upper bound on $\max(X_1, X_2, \dots, X_m)$ that holds with high probability over the random choice of the hashing function.

Bounding the max

If $\max(X_1, X_2, \dots, X_m) \geq l$ then there must be i such that $X_i \geq l$

$$P(\max(X_1, X_2, \dots, X_m) \geq l) \leq P(X_1 \geq l) + P(X_2 \geq l) + \dots + P(X_m \geq l) = mP(X_1 \geq l)$$

Using the union bound.

Suppose n (the number of items) grows to infinity, $m = \frac{n}{c}$ and $l(n)$ is a function of n
 c is fixed


Our goal, when computing an upper bound on the maximal occupancy is to

- A. Find the fastest increasing $l(n)$ such that $P(X_1 \geq l(n))$ decreases to zero with n .
- B. Find the slowest increasing $l(n)$ such that $mP(X_1 \geq l(n))$ decreases to zero with n .
- C. Find the slowest increasing $l(n)$ such that $P(X_1 \geq l(n))$ decreases to zero with n .

Bounding the deviation from the mean

- We know that the mean occupancy is n/m
- We want to show that the probability of a much higher occupancy is small.
- To do this we need to upper bound the deviation of the actual occupancy from the mean.
- We will do this using a sequence of better and better bounds, starting with Markov, then Chebyshev, and finally using the binomial coefficient.

Bounding using Markov inequality



Using Markov

$$E(X_i) = n / m = 1 / c$$

$$P(X_i \geq l(n)) \leq \frac{E(X_i)}{l(n)} = \frac{1}{cl(n)}$$

To get that $m(n)P(X_i \geq l(n)) \rightarrow 0$ we need that $\frac{m(n)}{cl(n)} \rightarrow 0$

$$m(n) = cn, \text{ so } \frac{n}{l(n)} \rightarrow 0$$


In other words, $l(n)$ increases faster than n .

But the max occupancy of any bin is n - the number of items

Using Markov generates a trivial bound.

Worst case bound using variance

Chebyshev's bound: $P(|X_i - E(X_i)| \geq k\sigma) \leq \frac{1}{k^2}$



$$\text{var}(X_i) = E(X_i^2) - E(X_i)^2; \quad E(X_i) = 1 / c$$

To bound $E(X_i^2)$ we think of X_i as a sum:

$$X_i = X_{i1} + \cdots X_{in}; \quad X_{ij} = \begin{cases} 1 & \text{if item } j \text{ is in bin } i \\ 0 & \text{otherwise} \end{cases}$$

X_{ij}, X_{ik} are independent if $j \neq k$

If two random variables X, Y are independent

then $E(XY) = E(X)E(Y)$

$$E(X_{ij}) = P(X_{ij} = 1) = 1 / m$$

Which of the following is true?

A. $E(X_{ij}X_{ik}) = 1 / m^2$

B. $E(X_{ij}X_{ik}) = 1 / m$ if $j = k$; $1 / m^2$ otherwise

C. $E(X_{ij}X_{ik}) = 1 / m$

Breaking up the square into a sum

$$E(X_i^2) = E((X_{i1} + \cdots + X_{in})^2) = ?$$

- A. $n^2 E(X_{ij}^2)$ B. $E(X_{i1} + \cdots + X_{in})^2$
C. $nE(X_{i1}^2) + n(n-1)E(X_{ij}X_{ik})$ D. $nE(X_{ii})$

Bounding the variance

$$nE(X_{i1}^2) + n(n-1)E(X_{ij}X_{ik})$$

$$E(X_{i1}^2) = E(X_{i1}) = P(\text{item 1 in bin } i) = 1/m$$

$$E(X_{ij}X_{ik}) = P(\text{item } j \text{ in bin } i \text{ and item } k \text{ in bin } i) = \\ = P(\text{item } j \text{ in bin } i)P(\text{item } k \text{ in bin } i) = 1/m^2$$

$$nE(X_{i1}^2) + n(n-1)E(X_{ij}X_{ik}) \leq \frac{n}{m} + \frac{n(n-1)}{m^2} \leq \frac{n}{m} + \frac{n^2}{m^2}$$

$$\text{if } m = cn, c \geq 1 \text{ we get } \text{var}(X_i) \leq \frac{1}{c} + \frac{1}{c^2} \leq \frac{2}{c}; \quad \sigma \leq \sqrt{\frac{2}{c}}$$

$$mP(X_i \geq l\sigma) \leq \frac{m}{l^2}$$

What is the slowest rate

that l can increase so that $(m / l^2) \rightarrow 0$ as $m \rightarrow \infty$?

A. $l = O(m^\alpha), \alpha > 1$ B. $l = O(\log m)$

C. $l = O(m^\alpha), \alpha > 1/2$ C. $l = O(m^\alpha), \alpha > 1/4$

A tighter bound using binomial coefficients

Instead of considering the mean and variance of X_i ,

We upper bound the probability using the binomial coefficient:

$$P(X_i \geq l) = \sum_{i=l}^n \binom{n}{i} \left(\frac{1}{m}\right)^i \left(\frac{m-1}{m}\right)^{n-i} \leq \binom{n}{l} \left(\frac{1}{m}\right)^l$$

Because:

We can first select the set of l that falls in the bin i and then select the remaining $n - l$ in an arbitrary way.

Standard inequality (see cheat-Sheets) : $\binom{n}{l} \leq \left(\frac{ne}{l}\right)^l$

$$\binom{n}{l} \left(\frac{1}{m}\right)^l \leq \left(\frac{ne}{l}\right)^l \left(\frac{1}{m}\right)^l = \left(\frac{ne}{lm}\right)^l = \left(\frac{e}{lc}\right)^l$$

We need $n \left(\frac{e}{lc}\right)^l \xrightarrow{n \rightarrow \infty} 0$

How fast does l need to grow in order to guarantee this?

A. $l = \log n$ B. $l = \log^2 n$ C. $l = \log \log n$ D. $l = \sqrt{n}$

The power of two choices

- We have shown that, if the ratio between items n and bins m is a constant $c = \frac{n}{m}$, and if $n \rightarrow \infty, m \rightarrow \infty$ then the probability that the highest occupancy is larger than $O(\log n)$ goes to zero.
- Is this the best that can be done?
- No!
- Mitzenmacher 1996, proposed the following ingenious method:
 - Instead of one hash function, use two (two different indices).
 - Given a new item a , compute both hash functions: $H_1(a), H_2(a)$
 - Compare the length of the lists in bin $H_1(a)$ and bin $H_2(a)$, add a to the end of the shorter list.
 - When retrieving search both lists for item a .
- Performance: maximal occupancy is $O(\log \log n)$ instead of $O(\log n)$

A general bound for binomial tails

Let X_1, X_2, \dots, X_n be independent, identically distributed

binary variables: $X_i = \begin{cases} 1 & \text{with prob. } p \\ 0 & \text{with prob. } 1-p \end{cases}$

Let $Y = \frac{1}{n}(X_1 + X_2 + \dots + X_n)$; $E(Y) = p$; $\text{var}(Y) = \frac{p(1-p)}{n}$

Then $P(Y \leq k) = \sum_{i=0}^k \binom{n}{i} p^i (1-p)^{n-i}$ The binomial tail.

We can bound the binomial tail using any of the following

for any $k > 1$, $P(Y \geq kp) \leq e^{-\frac{1}{3}np(k-1)^2}$

for any $k > 1$, $P\left(Y \leq \frac{p}{k}\right) \leq e^{-\frac{1}{2}np(1-1/k)^2}$

for any $\epsilon > 0$, $P(|Y - p| \geq \epsilon) \leq 2e^{-2n\epsilon^2}$