

Min-Hash and Document comparison

Yoav Freund

This lecture is based on:

- “Mining Massive Datasets” by Jure Leskovec, Anand Rajaraman and Jeffrey D. Ullman
- Description is in section 11.6 in the lecture notes.

Finding Similar Items

- ① Based on chapter 3 of the book “Mining Massive Datasets” by Jure Leskovec, Anand Rajaraman and Jeffrey D. Ullman
- ② Suppose we receive a stream of documents.
- ③ We want to find sets of documents that are very similar
- ④ Reasons: Plagiarism, Mirror web sites, articles with a common source.

Measuring the distance between sets

- 1 Suppose we consider the **set** of words in a document. Ignoring order and number of occurrences.
- 2 We will soon extend this assumption.
- 3 If two documents define two sets S, T , how do we measure the similarity between the two sets?
- 4 Jaccard similarity: $\frac{|S \cap T|}{|S \cup T|}$




Figure 3.1: Two sets with Jaccard similarity 3/8

Hash Functions

- ① Let X be a finite (but large) set
- ② Let $N = \{1, 2, \dots, n\}$ be a (very large) set of numbers.
- ③ A Hash-Function $h : X \rightarrow N$ is a function that “can be seen as” a mapping from each element of X to a an independently and uniformly chosen random element of N .
- ④ In actuality, the functions are pseudo-random and not random. But that distinction is out of scope for this class.

Min-Hash

- 1 Choose a random hash function h_i
- 2 Given a set of elements S in the domain X
- 3 $\text{min-}H_i(S) = \min_{s \in S} h_i(s)$
- 4 A min-hash **signature** for a document is the vector of numbers $\langle \text{min-}H_1(S), \text{min-}H_2(S), \dots, \text{min-}H_k(S) \rangle$
- 5 Signature also called a “sketch”: Any length document is represented by k numbers.
- 6 A lot of information is lost, but enough is retained to approximate the Jaccard similarity.

Visualizing Min-Hash

- We can represent the set of words in each document as a matrix.
- Rows a, b, c, \dots correspond to words.
- Columns S_1, S_2, \dots correspond to documents.
- A “1” in row b , column S_i means that document S_i contains the word b
- Hashing corresponds to randomly permuting the rows.
- **Min**-hashing a document corresponds to identifying the first “1” starting from the top of the column

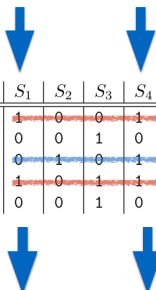
<i>Element</i>	S_1	S_2	S_3	S_4
a	1	0	0	1
b	0	0	1	0
c	0	1	0	1
d	1	0	1	1
e	0	0	1	0

<i>Element</i>	S_1	S_2	S_3	S_4
b	0	0	1	0
e	0	0	1	0
a	1	0	0	1
d	1	0	1	1
c	0	1	0	1

Understanding Min-Hash

- For any set S of size $|S|$, the probability that any particular element $s \in S$ is the min-hash is $1/|S|$
- Fix two documents S_i, S_j (columns) and partition the rows that contain at least a single “1” in those columns
- Denote by X rows that contain 1,1 (both documents contain the word.)
- Denote by Y rows that contain 1,0 or 0,1 (only one document contains the word)
- Permuting the rows does not change which rows are X and which are Y
- The min-hash of S_i, S_j agree if and only if first row that is not 0,0 is an X

- The probability that the min-hash of S_i, S_j agree is exactly $\frac{\#X}{\#X + \#Y}$
which is equal to $JS(S_i, S_j) = \frac{|S_i \cap S_j|}{|S_i \cup S_j|}$



Element	S_1	S_2	S_3	S_4	
a	1	0	0	1	X
b	0	0	1	0	
c	0	1	0	1	Y
d	1	0	1	1	X
e	0	0	1	0	

Element	S_1	S_2	S_3	S_4	
b	0	0	1	0	
e	0	0	1	0	
a	1	0	0	1	X
d	1	0	1	1	X
c	0	1	0	1	Y

Estimating Jaccard Similarity

- ① We can use min-hash to estimate Jaccard similarity (JS): $\frac{|S_i \cap S_j|}{|S_i \cup S_j|}$
- ② For each min hash function MH_i we have that

$$P_i [\text{min-}H_i(S) = \text{min-}H_i(T)] = \frac{|S \cap T|}{|S \cup T|}$$

- ③ A single comparison yields only true (1) or false (0)
- ④ Taking the average of k independent hash functions we can get an accurate estimate.

How many hash functions do we need? (1)

- ① From a statistics point of view we have k independent binary random variables:

$$X_i = \begin{cases} 1 & \text{if } \min\text{-}H_i(S) = \min\text{-}H_i(T) \\ 0 & \text{otherwise} \end{cases}$$

- ② We seek the expected value: $p \doteq E(X_i) = \frac{|S \cap T|}{|S \cup T|}$
- ③ We have to overcome the large std: $\sigma(X_i) = \sqrt{p(1-p)}$
- ④ Averaging gives a random variable with the same expected value but a smaller variance.

$$Y = \frac{1}{k} \sum_{i=1}^k X_i; \quad E(Y) = p \quad \sigma(Y) = \sqrt{\frac{p(1-p)}{k}}$$

⑤

$$\sigma(Y) \leq \sqrt{1/2(1-1/2)(1/k)} = \frac{1}{2\sqrt{k}}$$

Using a z-Scores to calculate the minimal number of hash functions.

- 1 Suppose we want our estimate of JS to be within ± 0.05 of the Jaccard distance with probability at least 95%
- 2 The fraction of min-has matches is the average of k independent binary random variables.
- 3 Lets assume k is large enough so that central limit theorem holds.
- 4 We want a confidence of 95% that the estimate is within ± 0.05 of the true value. In other words, we want

$$2\sigma(Y) \leq 0.05$$

- 5 Using the bound

$$\sigma(Y) \leq \frac{1}{2\sqrt{k}}$$

we find that it is enough if $\frac{1}{k} \leq 0.05$ or if $k \geq 20$

Introducing Order

- ❶ So far, we represented each document by the set of words it contains
- ❷ This removes the order in which the words appear: “Sampras beat Nadal” is the same as “Nadal beat Sampras”
- ❸ We can add order information to the set representation using **Shingles**

Shingles

- ① Consider the sentence: "the little dog loughed to see such craft"
- ② Word set representation: { "the", "little", "dog", "loughed", "to", "see", "such", "craft" }
- ③ 2-shingle representation: { "the little", "little dog", "dog loughed", "loughed to", "to see", "see such", "such craft" }
- ④ 3-shingle representation: { "the little dog", "little dog loughed", ... }
- ⑤ And so on
- ⑥ The number of shingles of length k from a document of length n is?
- ⑦ $n + 1 - k$ - largest for single words!
- ⑧ On the other hand, there is a much larger number of **different items**.
- ⑨ k too small - documents judged similar too often.
- ⑩ k too large - documents judged dissimilar too often

The problem with computing all pairwise distances

- 1 Suppose we have 1,000,000 documents, each containing about 10,000 words
- 2 Loading all documents to memory is infeasible
- 3 Using a min-hash signature for each document reduces storage to about 50 numbers per document. We can load into memory
- 4 Still, there are about $C(1,000,000, 2) \approx 10^{12}$ pairs
- 5 Computing *all* pairwise distances is infeasible
- 6 How about finding just the most similar pairs?

Finding only the similar pairs

- ① Idea 1: If the distance between **similar** pairs is much smaller than the average distance between random pairs then we can eliminate most pairs even with a crude approximation of the distance.
- ② Idea 2: Make a single pass over elements and map that to a **hash table** in such a way that distant pairs tend to fall in different bins.
- ③ Pairs that fall in the same bin are **"candidates"** - only candidates are compared to see whether they are actually similar pairs.
- ④ The process is repeated several times to increase the chance that all similar pairs become candidates.

Algorithm for efficiently finding the most similar pairs

- 1 Suppose we are using m functions for performing min-hash

- 2 Partition the functions to b bands, each containing r rows.
 $m = br$. In the figure
 $m = 12, b = 4, r = 3$

band 1	h1	1 0 0 0 2			
	h2	...	3 2 1 2 2	...	
	h3	0 1 3 1 1			
band 2	h4				
	h5				
	h6				
band 3	h7				
	h8				
	h9				
band 4	h10				
	h11				
	h12				


- 1 Use hash function to map the r -vector in each band to a hash table with K bins.
- 2 Candidate pairs are document that fall in the same bin in at least one of the bands.
- 3 Compute the similarity of the candidate pairs exactly
- 4 Accept those pairs whose similarity is higher than some threshold θ

The effect of b and r (1)

- ① Fix two documents (=two columns in the table). Suppose that the Jaccard similarity between them is s .
- ② The probability that the r -vectors are identical, is s^r , in this case the two r -vectors will map to the same bin in the hash table.
- ③ The probability that the the two r -vectors differ in at least on place is $1 - s^r$
- ④ The probability that all b signatures differ between the two documents is $(1 - s^r)^b$

The effect of b and r (2)

- 1 The probability that all b signatures differ between the two documents is $(1 - s^r)^b$
- 2 For $b = 20$, $r = 5$ we get the table on the right.



- For high similarity $s \geq 0.8$ the document pair is very likely to be a candidate.
- For low similarity $s \leq 0.2$ the document pair is very likely not to be a candidate.
- If 99.9% of the documents have low similarity then the number of false positives is about $10^{12} \times (.006 + 0.001) = 7 \times 10^9$. Still large, but not prohibitively so.

h a p p y

t h a n k s

g i v i n g